

INVESTOR DECK · CONFIDENTIAL

MODELS BUILT. COMPUTE NEXT.

A UK sovereign AI company with its MOTHER model family already trained and publicly released. This seed round funds the Dundee data centre that unlocks full product deployment at scale.

MOTHER CORE

MOTHER EXO

MSAI DUNDEE

AI FACTORY CLOUD

INTUITIV

THE OPPORTUNITY

THE HARD PART IS DONE. THE RAISE UNLOCKS **THE REST**

Most seed-stage AI companies raise to prove a model can be built. We have already built ours, and released it publicly. This round funds the sovereign compute that turns proven models into deployed revenue.

MODELS • BUILT & RELEASED

MOTHER CORE V3 is live and public on Hugging Face, benchmarked and downloadable. EXO, DEFENCE and T2V are built and ready.

INFRASTRUCTURE • ENGINEERED

MSAI Dundee fully designed: free-cooling, power and fit-out modelled and ready to build. Construction is what this raise funds.

PRODUCTS • READY TO SCALE

IntuiTV platform is live with T2V in demo; full product release is gated on the compute Dundee will provide.

V3

MODEL LIVE ON HUGGING FACE

77%

AGENTIC BENCHMARK • 81/105

~\$290k

PARTNER COMPUTE SECURED

£4M + \$66.3M

ASSET FINANCE • DRAWN +
EPOKA COVER

WE BUILT A MODEL FAMILY ON ~\$290K OF PARTNER COMPUTE

Before raising a seed round, MSAI trained its complete MOTHER family of sovereign models using non-dilutive compute credits from two programmes — demonstrating unusual capital efficiency for a full-stack AI company.

● NVIDIA INCEPTION

\$100,000 compute credit. Used for MOTHER CORE reasoning training and DeepVision experimentation on Blackwell-class hardware.

● LENOVO AI INNOVATORS

~\$190,000 of compute: a 3-month H200 8-GPU node running 24/7. Applied to MOTHER EXO embodied-AI training and multi-head vision pipeline development.

The point for investors: the hard, expensive work of building sovereign models is already done and paid for with partner credits. Seed capital compounds an existing asset rather than funding a science project.

THE **MOTHER** FAMILY OF SOVEREIGN MODELS

A UK-owned model stack spanning reasoning, embodied intelligence, defence and generative media — trained on sovereign infrastructure with full weight ownership.

MOTHER CORE • LIVE ON HF

Sovereign reasoning + agentic model. Public on Hugging Face, benchmarked at 77% on our agentic suite. Maths, science, tool-use and UK knowledge.

MOTHER EXO • BUILT • SCALING

Embodied / robotics intelligence for object, situational and historical awareness, targeting a humanoid platform. Full training scales with Dundee.

MOTHER DEFENCE • BUILT • SCALING

Special project: reasoning and strategy model for defence, aligned to recognised red-teaming and safety frameworks. Compute-gated for full release.

MOTHER T2V / DEEPVISION • DEMO

Sovereign text-to-video and vision stack powering IntuiTV. Live in limited demo; full broadcast-grade release scales with compute.

SOVEREIGN WEIGHTS

UK-HOSTED

FULL OWNERSHIP

MOTHERAI.UK/MODELS

LIVE & PUBLIC ON HUGGING FACE

MOTHER CORE V3 is not a slide claim. It is a released, downloadable, independently verifiable sovereign model — proof the MSAI pipeline produces real, working AI.

- ▶ ~7B sovereign reasoning + agentic model, custom transformer trained in-house.
- ▶ Public on Hugging Face: MediaStreamAI/MOTHER_CORE_V3, served via `trust_remote_code`.
- ▶ Purpose-built scope: reasoning, multi-step tool-use, RAG, and UK domain knowledge including Welsh, Irish and Scottish Gaelic.
- ▶ Full weight ownership: no foreign base-model dependency, no licence exposure.
- ▶ Training ongoing toward 80%+; Phase 2 adds coding and the EXO world-model backbone.

81/105

AGENTIC BENCHMARK • 77% (V2 WAS 49%)

7B

SOVEREIGN PARAMS • PUBLIC ON HF

100%

UK-OWNED WEIGHTS & PIPELINE

MOTHER EXO Sim

Sovereign UK Embodied AI · Fourier N1 — 1245×441×202 mm · 23-DOF · 1 kHz L0 · 475 Wh

telemetry stream error — sim backend not reachable

telemetry offline

idle

Reset

E-STOP

connecting...

MODELS · MOTHER EXO

EMBODIED AI FOR THE PHYSICAL WORLD

MOTHER EXO is a next-generation model for item, object, situational and historical awareness, built and in training for deployment inside a humanoid robotic shell. Full-scale training scales with Dundee compute.

MULTI-HEAD

Joint vision + reasoning heads for real-time scene understanding.

HUMANOID TARGET

Fourier N1 · 23-DOF, with MOTHER SPINE on-board compute.

COGNITIVE STACK

GR00T + DeepVision + CORE 7B, GUARDIAN-filtered actions.

BATTERY

0%

V: 0.00

P: 0 W

Runtime: 0h 0m

SIM METRICS

Sim time0.00 s

Wall time0.00 s

RTF0.00x

Ticks0

Dropped0

COGNITIVE STACK

L0Fourier N1 motor (1 kHz)0 Hz

L1Sim2Real ONNX0 policies

L2Isaac GR00T N1.6-MSAIstub

L2.5DEEPPVISION V2 scene graph0 obj

L3CORE 7B / ROBOTICS 7B-N1local

L4GUARDIAN signed action filterledger 0

GB10Blackwell SSH bridge-

ling-box and
ene
ion core behind

Powers MOTHER
DEFENCE &
OverWatch security
systems.

Chris ▾

MSAI Scotland (DUNDEE) : DC-3 Final

Netseed DC1-A · Scotland · Dundee · to scale (m)

544 GPU racks 688 racks 48x35 m 4 MW IT

WHITE SPACE · 48 × 35 m

DISCIPLINES

- Structure
- White Space
- Power
- Cooling
- Fibre / MMR
- Plant

INFRASTRUCTURE · MSAI DUNDEE

ENGINEERED. READY TO BUILD.

MSAI Dundee will be our first data centre — a sovereign GPU facility at Riverside, fully designed and modelled across power, cooling, white space and fibre. This raise funds construction. The images here are our engineering CAD and simulation, not yet a live hall.

544

GPU RACKS · DESIGNED

48x35m

WHITE SPACE HALL

10 → 25MW

PLANNED RAMP CAPACITY

Build

PRIMARY USE OF THIS RAISE

DESIGNED FOR PUE 1.001

HORIZON is our closed-loop free-cooling design: IT heat rejected via glycol primary and titanium plate heat exchanger into an on-site storm well — no river dependency. Figures below are from our simulation model.

100%

FREE COOLING DUTY
(MODELLED)

1.001

MODELLED PUE

13.3°C

STORM-WELL SUPPLY
(SIM)



FREE COOLING · 100% free

Pause

Reset

+120s/s

Storm well

13.3°C

Free cooling

100%

Primary supply

24.0°C

Primary return

32.8°C

Sec. flow

150 L/s

Prim. flow

120 L/s

Chiller duty

0.00 MW

PUE

1.001

Dundee air

4.2°C

Makeup

2.3 L/s

OPERATING POINT

IT load

4MW

Primary flow

120L/s

Secondary flow

150L/s

Well volume

320m³

Well recharge

650kW/K

PHE effect.

0.86

Rack supply

24°C

Chiller COP

3.5

Month

Jan · 4.2°C

Closed-loop: IT heat + glycol primary + titanium PHE + well-water secondary + onsite storm well (groundwater sink). Free cooling holds while the well + PHE approach meet the rack supply; otherwise the air-cooled chiller trims the balance.

PRODUCT • INTUITV

PLATFORM LIVE. SCALE PENDING COMPUTE.

• UK Sovereign AI • MOTHER T2V Now Live

IntuiTV is our live AI-native media platform — creation studio and broadcaster-grade HLS playout are operational. MOTHER T2V runs in limited demo today; full release scales once Dundee compute is online.

MOTHER T2V: A sovereign, UK-hosted text-to-video system for **broadcast-grade content generation**. Script to screen in minutes.

LIVE NOW

Studio + playout platform operational; creation tools in use.

T2V • DEMO

MOTHER text-to-video in limited demo, capacity-constrained.

COMPUTE-GATED

Full multi-channel release unlocks with Dundee GPU capacity.

[Watch Demo](#)

SOVEREIGN GPU-AS-A-SERVICE

gpu.mediastreamai.com — sovereign GPU capacity from MSAI Dundee, our first data centre. Capacity comes online as Dundee is built; the pipeline below is already contracted against it via LOIs and proposals.

BLACKWELL B300

Flagship training & inference clusters, liquid-cooled, RoCE fabric — spec'd and quoted.

H200 NVL

Inference capacity for production workloads, deploying with DC1.

COLOCATION

Single-tenant halls from 10MW, gas-primary power, ramping to 25MW.

COMMERCIAL PIPELINE — SIGNED LOIS & PROPOSALS (PIPELINE AGAINST FUTURE CAPACITY, NOT RAISED CAPITAL)

COUNTERPARTY	PRODUCT	STRUCTURE	INDICATIVE TCV
Cognizant	GPUaaS · 5,000 B300	LOI · 48-month	~£788M
Canopy Cloud	GPUaaS · 2,000 H200	LOI · 24-month	~£88M
Civo	Colocation · 10 – 25MW	Proposal delivered	Multi-year
Sesterce	DC3 Colocation	LOI · non-binding	Multi-year

CAPITAL & COMPUTE SECURED

What is confirmed and in place today — before this seed round. Kept deliberately separate from forecast and commercial pipeline.

SOURCE	TYPE	STATUS	VALUE
Nvidia Inception	Compute credit	Secured	\$100k
Lenovo AI Innovators	Compute • H200 3mo 8-GPU 24/7	Secured	~\$190k
Greymax Capital	Asset finance • drawn	Secured	£4.0M
EPOKA — B300 RVI	Asset finance • residual cover	Signed	\$37.52M
EPOKA — H200 RVI	Asset finance • residual cover	Signed	\$28.76M

~\$290k

NON-DILUTIVE COMPUTE

£4.0M + \$66.3M

ASSET FINANCE • GREYMAX DRAWN + EPOKA
RESIDUAL COVER

Signed

B300 + H200 RVI • GPU DOWNSIDE PROTECTED

THE RAISE BUILDS THE COMPUTE

Proceeds are weighted to the one thing that unlocks everything else: the Dundee data centre and the GPU compute that lets our built models and products deploy at full scale.



Use-of-funds split is indicative and will be finalised with lead investors. The senior GPUaaS debt facility (in arrangement) is expected to fund the majority of GPU hardware, with equity focused on infrastructure and scaling.

A CLEAR PATH TO SCALE REVENUE

Independent KPMG-built financial model. These are **forecast projections**, driven by GPUaaS, MOTHER model services and colocation rental — not capital raised.



£241M

FORECAST REVENUE • YR 5

~62%

FORECAST EBITDA MARGIN AT SCALE

KPMG

INDEPENDENT MODEL • JUNE 2026

Forecast only. Not a guarantee of future performance. Full model available in the data room.

WHY **MSAI** WINS

FULL-STACK SOVEREIGNTY

We own the models, the weights, and the products — and Dundee will give us the data centre to run them. Very few companies control the entire stack, fewer still in the UK.

PROOF, NOT PROMISE

The model is public and benchmarked on Hugging Face today. Seed capital scales a proven asset, not an unproven idea.

COMPUTE ADVANTAGE

Owning Dundee once built means margin control and no cloud-rent tax on every workload — unlike model-only competitors.

DEFENCE & SOVEREIGN DEMAND

MOTHER DEFENCE and UK sovereign positioning open procurement channels closed to foreign-model vendors.

Vertical integration by design: once Dundee is built, our models run on our GPUs in our own data centre, sold through our own products — each layer a customer for the next.

Drop anything in the

Open Studio

Sign in

JOIN THE ROUND

BACK THE UK'S SOVEREIGN AI FACTORY

Sign in to see your shared content.

Models built and released. Infrastructure engineered. Products live and ready to scale. This seed round funds the compute that turns proof into production.

Open Studio to publish content

Sign in